

Amazon Fine Food Reviews Analysis- sentimental analysis using Naive Bayes

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10 Attribute Information: 1) Id 2) ProductId - unique identifier for the product 3) UserId - unique identifier for the user 4) ProfileName 5) HelpfulnessNumerator - number of users who found the review helpful 6) HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not 7) Score - rating between 1 and 5 8) Time - timestamp for the review 9) Summary - brief summary of the review 10) Text - text of the review Objective:

Given a review, determine whether the review is positive (Rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use the Score/Rating. A rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

Loading the data

The dataset is available in two forms 1).csv file 2)SQLite Database In order to load the data, We have used the SQLITE dataset as it easier to query the data and visualise the data efficiently. Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

we will make positive label as 1 and negative score as 0 because we are taking log loss as metric.

Exploratory Data Analysis

In [4]:

```
#importing required libraries

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import sqlite3

import re
import nltk
import matplotlib.pyplot as plt
from nltk.corpus import stopwords
from nltk.stem import SnowballStemmer
import seaborn as sn
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import log_loss
from sklearn.metrics import accuracy_score
from sklearn.feature_extraction.text import CountVectorizer
```

In []:

```
#loading the data using sqlite
con = sqlite3.connect(r'C:\Users\welcome\Downloads\amazon-fine-food-reviews\database.sqlite')
data = pd.read_sql_query('select * from reviews where Score !=3',con)
data.head()
```

In []:

```
def change_labels(x):
    if int(x) > 3:
        return 1
    return 0
```

In []:

In []:

Text Preprocessing: Stemming, stop-word removal and Lemmatization.

Hence in the Preprocessing phase we do the following in the order below:-

- In [22]:

In []:

```
import datetime

str1=' '
final_string=[]

s=' '
```

```

start_time = datetime.datetime.now()
for sent in clean_data['Text'].values:
    filtered_sentence=[]
    sent=clean_html(sent) # remove HTML tags
    for w in sent.split():
        for cleaned_words in cleansen(w).split():
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                if(cleaned_words.lower() not in stop_words):
                    s=(stemmer.stem(cleaned_words.lower())).encode('utf8')
                    filtered_sentence.append(s)
                else:
                    continue
            else:
                continue
        #print(filtered_sentence)
    str1 = b" ".join(filtered_sentence) #final string of cleaned words

    final_string.append(str1)
clean_data['CleanedText']=final_string
print('Total time taken to clean the reviews',datetime.datetime.now()-start_time)

```

In []:

```
clean_data['CleanedText'].head()
```

In []:

```
final_data = clean_data.head(100000)
final_data.shape
```

In []:

```
final_data.to_pickle("C:/Users/welcome/Downloads/amazon-fine-food-reviews/final_data.pkl")
```

In [5]:

```
data = pd.read_pickle("C:/Users/welcome/Downloads/amazon-fine-food-reviews/final_data.pkl")
```

In [6]:

```
data.shape
```

Out[6]:

```
(100000, 11)
```

Split the dataset randomly into three parts train, cross validation and test with 64%,16%, 20% of data respectively

In [7]:

```

#noe lets split the data
from sklearn.model_selection import train_test_split
x,x_test,y,y_test = train_test_split(data['CleanedText'],data['Score'],train_size=0.8,shuffle=False)
x_train,x_cv,y_train,y_cv = train_test_split(x,y,train_size=0.8,shuffle=False)

```

J:\ANACONDA3\lib\site-packages\sklearn\model_selection_split.py:2026: FutureWarning: From version 0.21, test_size will always complement train_size unless both are specified.
FutureWarning)

Bag of Words (BoW)

A commonly used model in methods of Text Classification. As part of the BOW model, a piece of text (sentence or a document) is represented as a bag or multiset of words, disregarding grammar and even word order and the frequency or occurrence of each word is used as a feature for training a classifier.

In [8]:

```
#Lets Vecotirize
#bagof words
bag_words = CountVectorizer()
x_train_bag= bag_words.fit_transform(x_train)
x_test_bag= bag_words.transform(x_test)
x_cv_bag= bag_words.transform(x_cv)

print('After vectorizing shape of x Train',x_train_bag.shape)
print('After vectorizing shape of x Test',x_test_bag.shape)
print('After vectorizing shape of x CV',x_cv_bag.shape)
```

After vectorizing shape of x Train (64000, 30957)
After vectorizing shape of x Test (20000, 30957)
After vectorizing shape of x CV (16000, 30957)

Multinomial Naive bayees

MultinomialNB implements the naive Bayes algorithm for multinomially distributed data, and is one of the two classic naive Bayes variants used in text classification (where the data are typically represented as word vector counts)

In [9]:

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import log_loss,accuracy_score,confusion_matrix,f1_score
import seaborn as sns
from sklearn.metrics import classification_report
import matplotlib.ticker as ticker
import matplotlib.ticker as plticker

#lets do hyper parameter tuning now

range1 = [10 ** x for x in range(-4,4)]
range2 = [ 2 ** x for x in range(-4,4)]
#alpha = range1 + range2
alpha = [range1,range2]
for r in alpha:

    cv_f1_score_array=[]
    train_f1_score_array=[]
    for i in r:
        clf_bag = MultinomialNB(alpha = i )
        clf_bag.fit(x_train_bag,y_train)
        predict_y = clf_bag.predict(x_cv_bag)
        predict_y_train = clf_bag.predict(x_train_bag)
        cv_f1_score_array.append(f1_score(y_cv, predict_y))
        train_f1_score_array.append(f1_score(y_train, predict_y_train))
        print('For values of alpa = ', i, "The f1 score",f1_score(y_cv, predict_y))
    if(r == range1):
        fig, ax = plt.subplots(figsize=(20,7))

        plt.xscale('log')
        plt.plot(r, cv_f1_score_array,c='g')
        plt.plot(r, train_f1_score_array,c='b')
        for i, txt in enumerate(np.round(cv_f1_score_array,3)):
            ax.annotate((r[i],np.round(txt,3)), (r[i],cv_f1_score_array[i]))
        for i, txt in enumerate(np.round(train_f1_score_array,3)):
            ax.annotate((r[i],np.round(txt,3)), (r[i],train_f1_score_array[i]))
        plt.grid()
        plt.title("f1_score for each alpha ")
        plt.xlabel("Aplha values's")
        plt.ylabel("f1_score")
        plt.show()

    else:
        fig, ax = plt.subplots(figsize=(20,7))
        ax.set_xscale('log', basex=2)


        plt.plot(r, cv_f1_score_array,c='g')
        plt.plot(r, train_f1_score_array,c='b')
        for i, txt in enumerate(np.round(cv_f1_score_array,3)):
            ax.annotate((r[i],np.round(txt,3)), (r[i],cv_f1_score_array[i]))
        for i, txt in enumerate(np.round(train_f1_score_array,3)):
            ax.annotate((r[i],np.round(txt,3)), (r[i],train_f1_score_array[i]))
```

```

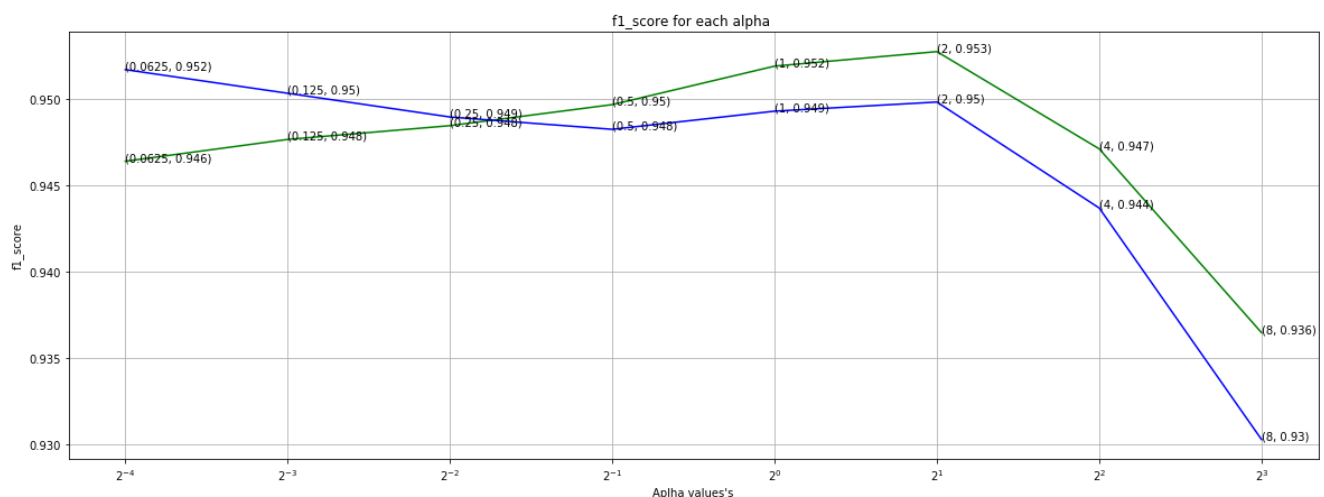
plt.grid()
plt.title("f1_score for each alpha ")
plt.xlabel("Alpha values's")
plt.ylabel("f1_score")
plt.show()

```

For values of alpha = 0.0001 The f1 score 0.9408101918121957
 For values of alpha = 0.001 The f1 score 0.9424463011439022
 For values of alpha = 0.01 The f1 score 0.9447402503957404
 For values of alpha = 0.1 The f1 score 0.9471745802491425
 For values of alpha = 1 The f1 score 0.9519179208522504
 For values of alpha = 10 The f1 score 0.9331940137521909
 For values of alpha = 100 The f1 score 0.9282604327148504
 For values of alpha = 1000 The f1 score 0.9282963260658428



For values of alpha = 0.0625 The f1 score 0.946403435954813
 For values of alpha = 0.125 The f1 score 0.9476687493228357
 For values of alpha = 0.25 The f1 score 0.9484662576687117
 For values of alpha = 0.5 The f1 score 0.9496868926797668
 For values of alpha = 1 The f1 score 0.9519179208522504
 For values of alpha = 2 The f1 score 0.9527600662274984
 For values of alpha = 4 The f1 score 0.9471215204718706
 For values of alpha = 8 The f1 score 0.9364654909792506



By looking aboce plots we can say alpha value as 0.25

In [14]:

```

best_alpha_bow = 0.25
clf_bag_best = MultinomialNB(alpha = best_alpha_bow )
clf_bag_best.fit(x_train_bag, y_train)
predict_y = clf_bag_best.predict(x_train_bag)
print('f1score on train data is :',f1_score(y_train,predict_y))
print('Accuracy on train data is :',accuracy_score(y_train,predict_y)*100)

```

```

predict_y = clf_bag_best.predict(x_test_bag)

acc_bag = accuracy_score(y_test,predict_y)
print('f1score on test data is :',f1_score(y_test,predict_y))
print('Accuracy on test data is ',acc_bag*100)

C = confusion_matrix(y_test, predict_y)

print("-"*20, "Confusion matrix", "-"*20)
plt.figure(figsize=(20,7))
plt.figure(figsize=(20,7))
sns.heatmap(C, annot=True, cmap="YlGnBu", fmt="d")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
print(classification_report(y_test, predict_y))


```

```

f1score on train data is : 0.9489689956979677
Accuracy on train data is : 91.4
f1score on test data is : 0.9341633318511932
Accuracy on test data is  88.895
----- Confusion matrix -----

```

<Figure size 1440x504 with 0 Axes>



	precision	recall	f1-score	support
0	0.65	0.64	0.65	3168
1	0.93	0.94	0.93	16832
avg / total	0.89	0.89	0.89	20000

In [12]:

```

#printing the important features
log_prob = clf_bag_best.feature_log_prob_
bag_features = bag_words.get_feature_names()
feature_prob = pd.DataFrame(log_prob, columns = bag_features)
feature_prob_bag = feature_prob.T
feature_prob_bag.shape

```

Out[12]:
(30957, 2)

In [13]:

```

print("Top 20 negative features:-\n",feature_prob_bag[0].sort_values(ascending = False)[0:20])
print("\n\n Top 20 positive features:-\n",feature_prob_bag[1].sort_values(ascending = False)[0:20])

```

```
Top 20 negative features:-
```

```
tast      -4.384116
like      -4.393441
product   -4.451957
one       -4.750718
would     -4.944157
tea       -4.969019
tri       -4.974960
flavor    -5.051556
use       -5.068587
food      -5.075151
good      -5.108638
get       -5.148587
buy       -5.169628
order     -5.196220
dont      -5.274474
dog       -5.287623
box       -5.307159
bag       -5.335593
even      -5.392744
eat       -5.409855
```

```
Name: 0, dtype: float64
```

```
Top 20 positive features:-
```

```
tea       -4.446648
like      -4.543737
tast      -4.629236
use       -4.670388
good      -4.708215
love      -4.710128
great     -4.743731
one       -4.813980
flavor    -4.832914
product   -4.916880
tri       -4.966461
make      -4.993692
get       -5.092413
food      -5.218467
time      -5.298046
find      -5.412826
buy       -5.415702
eat       -5.430659
best      -5.443335
dog       -5.454444
```

```
Name: 1, dtype: float64
```

Observations

1)From above plot(F1_score vs optimal alpha) for alpha =0.25 we are getting good f1 train and test scores 2)In confusion matrix, we can see diagonal elements values are high which means our model performing good 3)with unseen data,we are getting accuracy of 88.85%

TF-IDF

```
In [15]:
```

```
tfidf_words = TfidfVectorizer()
x_train_tfidf= tfidf_words.fit_transform(x_train)
x_test_tfidf= tfidf_words.transform(x_test)
x_cv_tfidf = tfidf_words.transform(x_cv)

print('After vectorizing shape of x Train',x_train_tfidf.shape)
print('After vectorizing shape of x Test',x_test_tfidf.shape)
print('After vectorizing shape of x CV',x_cv_tfidf.shape)
```

```
After vectorizing shape of x Train (64000, 30957)
```

```
After vectorizing shape of x Test (20000, 30957)
```

```
After vectorizing shape of x CV (16000, 30957)
```

In [16]:

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import log_loss, accuracy_score, confusion_matrix
import seaborn as sns
from sklearn.metrics import classification_report
#lets do hyper parameter tuning now
#lets do hyper parameter tuning now

range1 = [10 ** x for x in range(-4,4)]
range2 = [ 2 ** x for x in range(-6,6)]
#alpha = range1 + range2
alpha = [range1,range2]
for r in alpha:


    cv_f1_score_array=[]
    train_f1_score_array=[]
    for i in r:
        clf_bag = MultinomialNB(alpha = i )
        clf_bag.fit(x_train_tfidf,y_train)
        predict_y = clf_bag.predict(x_cv_tfidf)
        predict_y_train = clf_bag.predict(x_train_tfidf)
        cv_f1_score_array.append(f1_score(y_cv, predict_y))
        train_f1_score_array.append(f1_score(y_train, predict_y_train))
        print('For values of alpha = ', i, "The f1 score",f1_score(y_cv, predict_y))
    if(r == range1):
        fig, ax = plt.subplots(figsize=(20,7))


        plt.xscale('log')
        plt.plot(r, cv_f1_score_array,c='g')
        plt.plot(r, train_f1_score_array,c='b')
        for i, txt in enumerate(np.round(cv_f1_score_array,3)):
            ax.annotate((r[i],np.round(txt,3)), (r[i],cv_f1_score_array[i]))
        for i, txt in enumerate(np.round(train_f1_score_array,3)):
            ax.annotate((r[i],np.round(txt,3)), (r[i],train_f1_score_array[i]))
        plt.grid()
        plt.title("f1_score for each alpha ")
        plt.xlabel("Alpha values's")
        plt.ylabel("f1_score")
        plt.show()

    else:
        fig, ax = plt.subplots(figsize=(20,7))
        ax.set_xscale('log', base=2)


        plt.plot(r, cv_f1_score_array,c='g')
        plt.plot(r, train_f1_score_array,c='b')
        for i, txt in enumerate(np.round(cv_f1_score_array,3)):
            ax.annotate((r[i],np.round(txt,3)), (r[i],cv_f1_score_array[i]))
        for i, txt in enumerate(np.round(train_f1_score_array,3)):
            ax.annotate((r[i],np.round(txt,3)), (r[i],train_f1_score_array[i]))
        plt.grid()
        plt.title("f1_score for each alpha ")
        plt.xlabel("Alpha values's")
        plt.ylabel("f1_score")
        plt.show()
```

For values of alpha = 0.0001 The f1 score 0.9367279197892072
For values of alpha = 0.001 The f1 score 0.9372926003215765
For values of alpha = 0.01 The f1 score 0.9397005537704246
For values of alpha = 0.1 The f1 score 0.9402181322426721
For values of alpha = 1 The f1 score 0.9300359048354082
For values of alpha = 10 The f1 score 0.9282963260658428
For values of alpha = 100 The f1 score 0.9282963260658428
For values of alpha = 1000 The f1 score 0.9282963260658428





For values of alpha = 0.015625 The f1 score 0.940404592673592
 For values of alpha = 0.03125 The f1 score 0.9405375866662113
 For values of alpha = 0.0625 The f1 score 0.9407227928881001
 For values of alpha = 0.125 The f1 score 0.9400824334911606
 For values of alpha = 0.25 The f1 score 0.9374830485489558
 For values of alpha = 0.5 The f1 score 0.9339492228328669
 For values of alpha = 1 The f1 score 0.9300359048354082
 For values of alpha = 2 The f1 score 0.9285092127303183
 For values of alpha = 4 The f1 score 0.9282963260658428
 For values of alpha = 8 The f1 score 0.9282963260658428
 For values of alpha = 16 The f1 score 0.9282963260658428
 For values of alpha = 32 The f1 score 0.9282963260658428



taking best alpha value as 0.25

In [17]:

```
best_alpha_tfidf = 0.25
clf_tfidf_best = MultinomialNB(alpha = best_alpha_tfidf )
clf_tfidf_best.fit(x_train_tfidf, y_train)
predict_y = clf_tfidf_best.predict(x_train_tfidf)
print('f1score on train data is :',f1_score(y_train,predict_y))
print('Accuracy on train data is :',accuracy_score(y_train,predict_y)*100)

predict_y = clf_tfidf_best.predict(x_test_bag)

acc_tfidf = accuracy_score(y_test,predict_y)
print('f1score on test data is :',f1_score(y_test,predict_y))
print('Accuracy on test data is :',acc_tfidf*100)

C = confusion_matrix(y_test, predict_y)


print("-"*20, "Confusion matrix", "-"*20)
plt.figure(figsize=(20,7))
plt.figure(figsize=(20,7))
sns.heatmap(C, annot=True, cmap="YlGnBu", fmt="d")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
print(classification_report(y_test, predict_y))
```

f1score on train data is : 0.9383901811240574
 Accuracy on train data is : 88.8546875
 f1score on test data is : 0.9420493282534589

Accuracy on test data is 89.885

----- Confusion matrix -----

<Figure size 1440x504 with 0 Axes>



	precision	recall	f1-score	support
0	0.80	0.48	0.60	3168
1	0.91	0.98	0.94	16832
avg / total	0.89	0.90	0.89	20000

In [18]:

```
#printing the important features
log_prob = clf_tfidf_best.feature_log_prob_
tfidf_features = tfidf_words.get_feature_names()
feature_prob = pd.DataFrame(log_prob, columns = tfidf_features)
feature_prob_bag = feature_prob.T
feature_prob_bag.shape
```

Out[18]:

(30957, 2)

In [19]:

```
print("Top 20 negative features:-\n",feature_prob_bag[0].sort_values(ascending = False)[0:20])
print("\n\n Top 20 positive features:-\n",feature_prob_bag[1].sort_values(ascending = False)[0:20])
```

Top 20 negative features:-

tast	-5.075296
product	-5.137854
like	-5.177555
would	-5.492188
one	-5.514936
tea	-5.529719
order	-5.625002
buy	-5.651377
tri	-5.677593
flavor	-5.688325
box	-5.711517
dog	-5.767227
dont	-5.767873
disappoint	-5.807237
good	-5.823675
get	-5.828331
food	-5.836529
bag	-5.836772
use	-5.904593
even	-5.925115

```
even          -5.920445  
Name: 0, dtype: float64
```

```
Top 20 positive features:-  
tea          -4.869499  
great        -5.096738  
love         -5.110994  
good         -5.194902  
tast         -5.228223  
like         -5.256520  
use          -5.309518  
product      -5.339280  
flavor       -5.352095  
one          -5.467646  
make         -5.537003  
tri          -5.541485  
get          -5.641047  
best         -5.657591  
dog          -5.686775  
find         -5.703858  
buy          -5.713891  
food         -5.719795  
time         -5.735033  
order        -5.739257  
Name: 1, dtype: float64
```

Observations

1)From above plot(f1score vs optimal alpha) for alpha=0.25 , we are getting good f1 score 2)In confusion matrix, we can see diagonal elements values are high which means our model performing good 3)with unseen data,we are getting accuracy of 89.85%

Lets test model by adding one more fetures i.e., summary column to text

In [20]:

```
data['additionaltext'] = data['Text'].str.cat(data['Summary'])  
data['additionaltext'].head()
```

Out[20]:

```
138706    this witty little book makes my son laugh at l...  
138688    I grew up reading these Sendak books, and watc...  
138689    This is a fun way for children to learn their ...  
138690    This is a great little book to read aloud- it ...  
138691    This is a book of poetry about the months of t...  
Name: additionaltext, dtype: object
```

In [25]:

```
import datetime  
  
str1=' '  
final_string=[]  
  
s=' '  
start_time = datetime.datetime.now()  
for sent in data['additionaltext'].values:  
    filtered_sentence=[]  
    sent=clean_html(sent) # remove HTML tags  
    for w in sent.split():  
        for cleaned_words in cleansen(w).split():  
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):  
                if(cleaned_words.lower() not in stop_words):  
                    s=(stemmer.stem(cleaned_words.lower()).encode('utf8'))  
                    filtered_sentence.append(s)  
                else:  
                    continue  
            else:  
                continue  
    #print(filtered_sentence)  
    str1 = b" ".join(filtered_sentence) #final string of cleaned words
```

```

    final_string.append(str1)
data['additionaltext']=final_string
print('Total time taken to clean the reviews',datetime.datetime.now()-start_time)

```

Total time taken to clean the reviews 0:04:17.429540

In [26]:

```
data['additionaltext'].head()
```

Out[26]:

```

138706    b'witti littl book make son laugh loud recit c...
138688    b'grew read sendak book watch realli rosi movi...
138689    b'fun way children learn month year learn poem...
138690    b'great littl book read nice rhythm well good ...
138691    b'book poetri month year goe month cute littl ...
Name: additionaltext, dtype: object

```

In [27]:

```

#noe lets split the data
from sklearn.model_selection import train_test_split
x,x_test_add,y,y_test_add = train_test_split(data['additionaltext'],data['Score'],train_size=0.8)
x_train_add,x_cv_add,y_train_add,y_cv_add = train_test_split(x,y,train_size=0.8)

```

J:\ANACONDA3\lib\site-packages\sklearn\model_selection_split.py:2026: FutureWarning: From version 0.21, test_size will always complement train_size unless both are specified.
FutureWarning)

In [28]:

```

#bagof words
bag_words = CountVectorizer()
x_train_bag_add= bag_words.fit_transform(x_train_add)
x_test_bag_add= bag_words.transform(x_test_add)
x_cv_bag_add= bag_words.transform(x_cv_add)

print('After vectorizing shape of x Train',x_train_bag_add.shape)
print('After vectorizing shape of x Test',x_test_bag_add.shape)
print('After vectorizing shape of x CV',x_cv_bag_add.shape)

```

After vectorizing shape of x Train (64000, 43751)
After vectorizing shape of x Test (20000, 43751)
After vectorizing shape of x CV (16000, 43751)

In [28]:

```

#lets do hyper parameter tuning now

rangel = [10 ** x for x in range(-4,4)]
range2 = [ 2 ** x for x in range(-4,4)]
#alpha = rangel + range2
alpha = [rangel,range2]
for r in alpha:

    cv_f1_score_array=[]
    train_f1_score_array=[]
    for i in r:
        clf_bag = MultinomialNB(alpha = i )
        clf_bag.fit(x_train_bag_add,y_train_add)
        predict_y = clf_bag.predict(x_cv_bag_add)
        predict_y_train = clf_bag.predict(x_train_bag_add)
        cv_f1_score_array.append(f1_score(y_cv_add, predict_y))
        train_f1_score_array.append(f1_score(y_train_add, predict_y_train))
        print('For values of alpa = ', i, "The f1 score",f1_score(y_cv_add, predict_y))
    if(r == rangel):
        fig, ax = plt.subplots(figsize=(20,7))

        plt.xscale('log')

```

```


plt.plot(r, cv_f1_score_array,c='g')
plt.plot(r, train_f1_score_array,c='b')
for i, txt in enumerate(np.round(cv_f1_score_array,3)):
    ax.annotate((r[i],np.round(txt,3)), (r[i],cv_f1_score_array[i]))
for i, txt in enumerate(np.round(train_f1_score_array,3)):
    ax.annotate((r[i],np.round(txt,3)), (r[i],train_f1_score_array[i]))
plt.grid()
plt.title("f1_score for each alpha ")
plt.xlabel("Alpha values's")
plt.ylabel("f1_score")
plt.show()

else:
    fig, ax = plt.subplots(figsize=(20,7))
    ax.set_xscale('log', basex=2)


    plt.plot(r, cv_f1_score_array,c='g')
    plt.plot(r, train_f1_score_array,c='b')
    for i, txt in enumerate(np.round(cv_f1_score_array,3)):
        ax.annotate((r[i],np.round(txt,3)), (r[i],cv_f1_score_array[i]))
    for i, txt in enumerate(np.round(train_f1_score_array,3)):
        ax.annotate((r[i],np.round(txt,3)), (r[i],train_f1_score_array[i]))
    plt.grid()
    plt.title("f1_score for each alpha ")
    plt.xlabel("Alpha values's")
    plt.ylabel("f1_score")
    plt.show()


```

For values of alpha = 0.0001 The f1 score 0.943751143641354
 For values of alpha = 0.001 The f1 score 0.9449807304092492
 For values of alpha = 0.01 The f1 score 0.9462254775663439
 For values of alpha = 0.1 The f1 score 0.9465524878192825
 For values of alpha = 1 The f1 score 0.9501268575570859
 For values of alpha = 10 The f1 score 0.924072505505675
 For values of alpha = 100 The f1 score 0.9205545997368688
 For values of alpha = 1000 The f1 score 0.9205235460801512



For values of alpha = 0.0625 The f1 score 0.9464437308529879
 For values of alpha = 0.125 The f1 score 0.9470460164581719
 For values of alpha = 0.25 The f1 score 0.9468954850114163
 For values of alpha = 0.5 The f1 score 0.9480200739953843
 For values of alpha = 1 The f1 score 0.9501268575570859
 For values of alpha = 2 The f1 score 0.9497083511168017
 For values of alpha = 4 The f1 score 0.9388023993620193
 For values of alpha = 8 The f1 score 0.9267795227411788





In [29]:

```
best_alpha_bow = 10
clf_bag_best = clf_bag_best = MultinomialNB(alpha = best_alpha_bow )
clf_bag_best.fit(x_train_bag_add, y_train_add)
predict_y = clf_bag_best.predict(x_train_bag_add)

print('f1score on train data is :',f1_score(y_train_add,predict_y))
print('Accuracy on train data is :',accuracy_score(y_train_add,predict_y)*100)

predict_y = clf_bag_best.predict(x_test_bag_add)

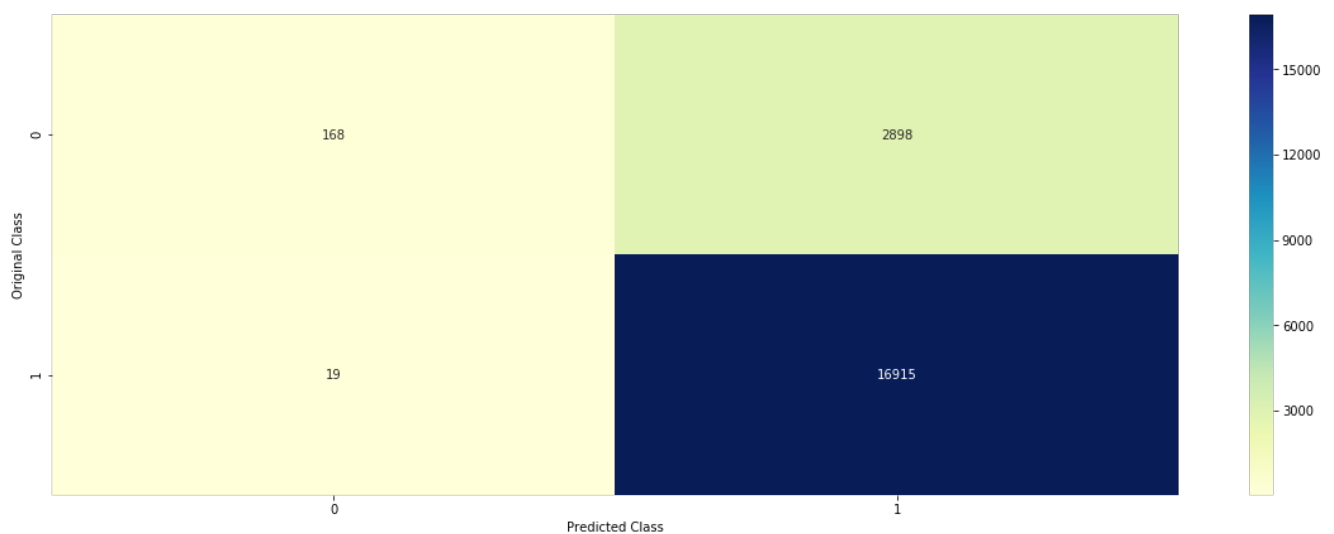
acc_tfidf = accuracy_score(y_test_add,predict_y)
print('f1score on test data is :',f1_score(y_test_add,predict_y))
print('Accuracy on test data is ',acc_tfidf*100)

C = confusion_matrix(y_test_add, predict_y)

print("-"*20, "Confusion matrix", "-"*20)
plt.figure(figsize=(20,7))
plt.figure(figsize=(20,7))
sns.heatmap(C, annot=True, cmap="YlGnBu", fmt="d")
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
print(classification_report(y_test, predict_y))
```

```
f1score on train data is : 0.9259155527265013
Accuracy on train data is : 86.3609375
f1score on test data is : 0.9206193702887311
Accuracy on test data is  85.41499999999999
----- Confusion matrix -----
```

<Figure size 1440x504 with 0 Axes>



	precision	recall	f1-score	support
0	0.18	0.01	0.02	3168
1	0.84	0.99	0.91	16832
avg / total	0.74	0.84	0.77	20000

Observations: Model is performing worst as we can see from confusion matrix its highly biased to majority class labels. one of the reason for bad performance is may be due to fact that summary and text columns are related . we know naive bayes assume that presence of a particular feature in a class is unrelated to presence of any other feature, this is the reason we will not apply naive bayes on word2vec because it tries to give relation between features.

Performance Table

□

Conclusion: Looking above table both of the models are performing well on unseen data.